

CPSC 481

Artificial Intelligence

Dr. Mira Kim

Mira.kim@fullerton.edu

What we will cover today




- Informed search
 - Heuristics
 - Greedy best-first search
 - A^*

Textbook: Chapter 3.5

Search Heuristics

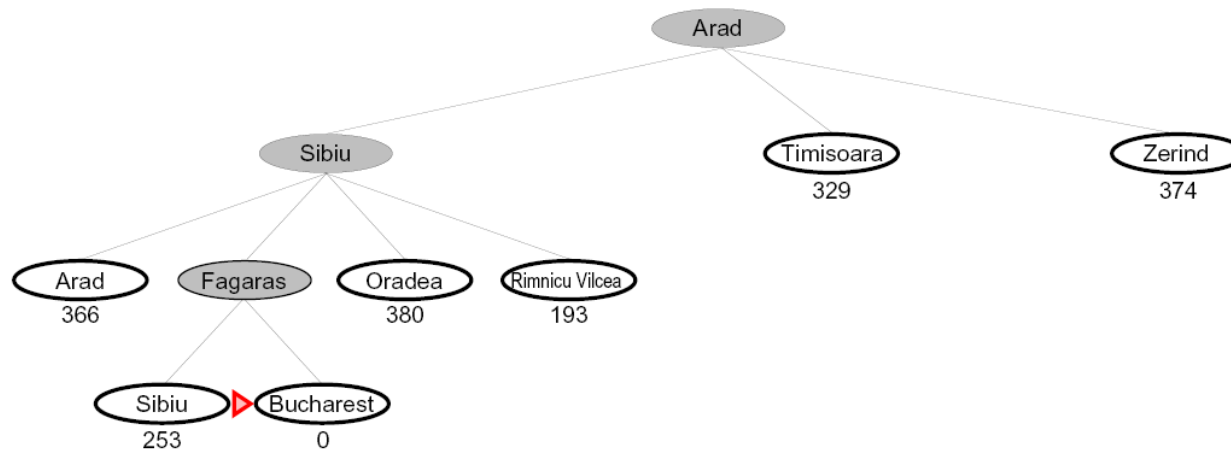


- A heuristic is:
 - A function that *estimates* how close a state is to a goal
 - Designed for a particular search problem





Greedy Best-First Search

- Expand the node that seems closest according to heuristic
- Heuristic: estimate of distance to nearest goal for each state
- Also called “Pure heuristic search”



Greedy Best-first search


- What can go wrong?
- Not optimal
- Can get stuck in dead ends
 - Not complete



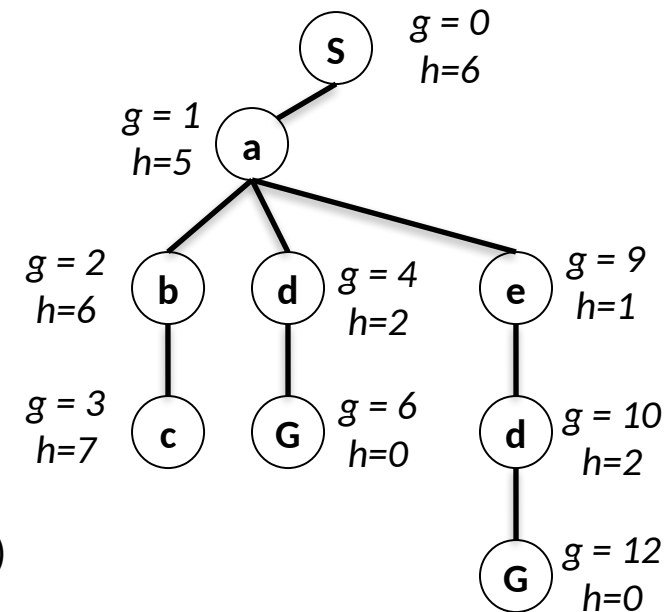
Heuristic cost
(perfectly accurate!)
Actual cost

A* Search – Combining UCS and Greedy

- **Uniform-cost** orders by path cost, or *backward cost* $g(n)$
- **Greedy** orders by goal proximity, or *forward cost* $h(n)$



- **A* Search** orders by the sum: $f(n) = g(n) + h(n)$
- A* is pronounced “A star”



Example: Teg Grenager

A* search

- Input: A problem
- Data structures:
 - Frontier (also called “open”)
 - Priority Queue, **ordered by path cost + heuristic cost**
 - Explored (also called “closed”)
 - Set, for efficiency

Initialize frontier with initial state

Initialize explored to empty

Loop do

IF the frontier is empty RETURN FAILURE

Choose lowest cost node from frontier and remove it

IF lowest cost node is goal RETURN SUCCESS

Add node to explored

FOR every child node of node

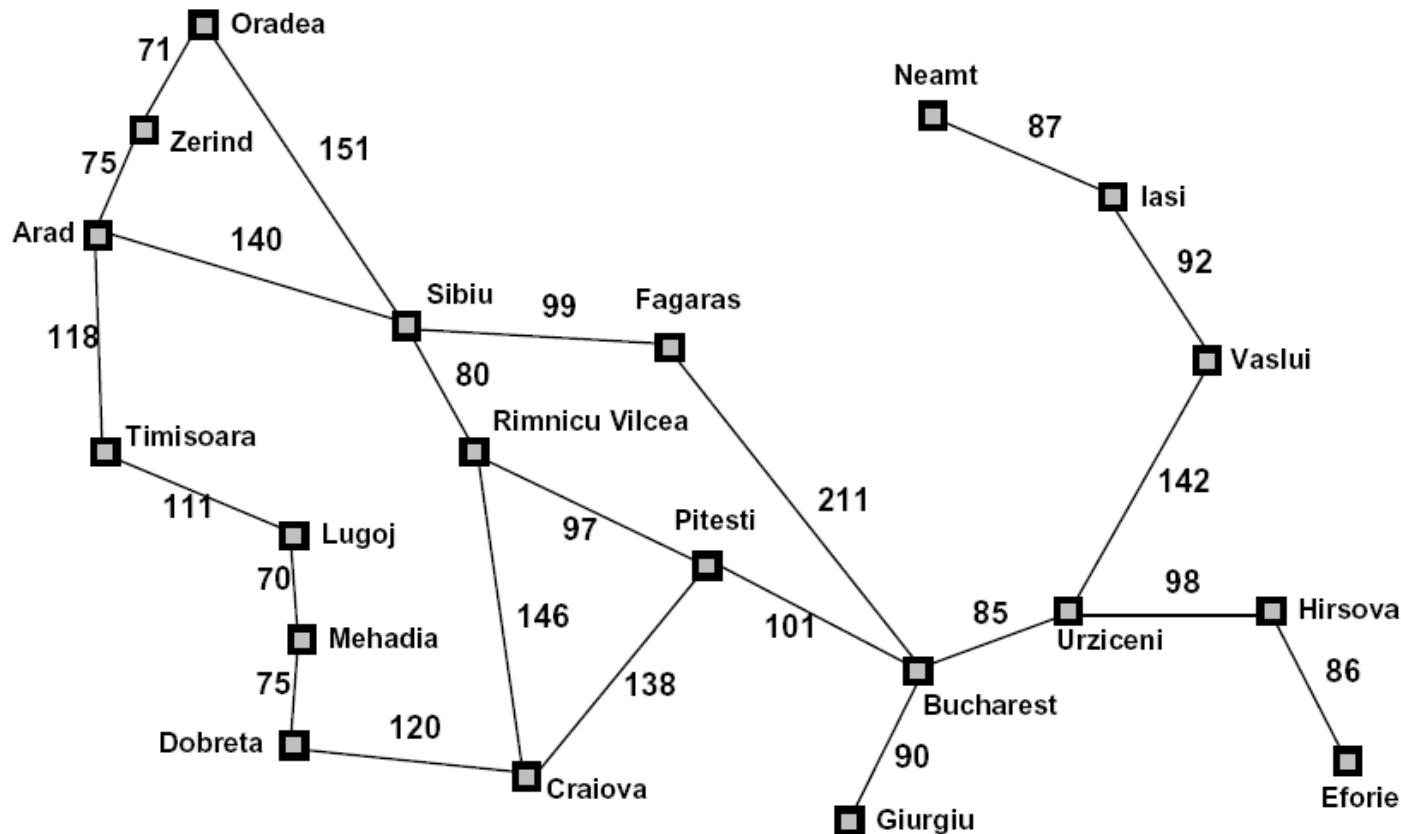
IF child not already on frontier or explored

insert child node to frontier

ELSE IF child is in frontier with higher path cost

replace existing frontier node with child node

Example: Heuristic Function

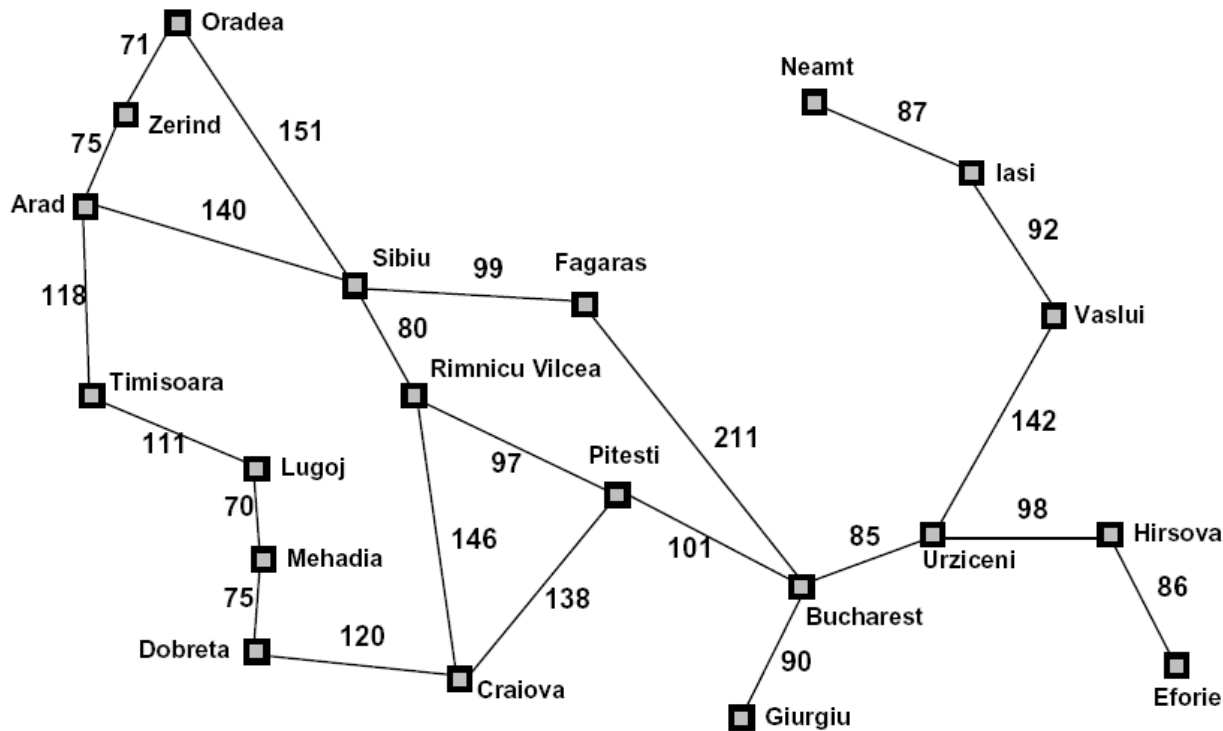


Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

$h(n)$

Explored (set)	Frontier (priority queue)
A(0+366=366)	{}




Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374



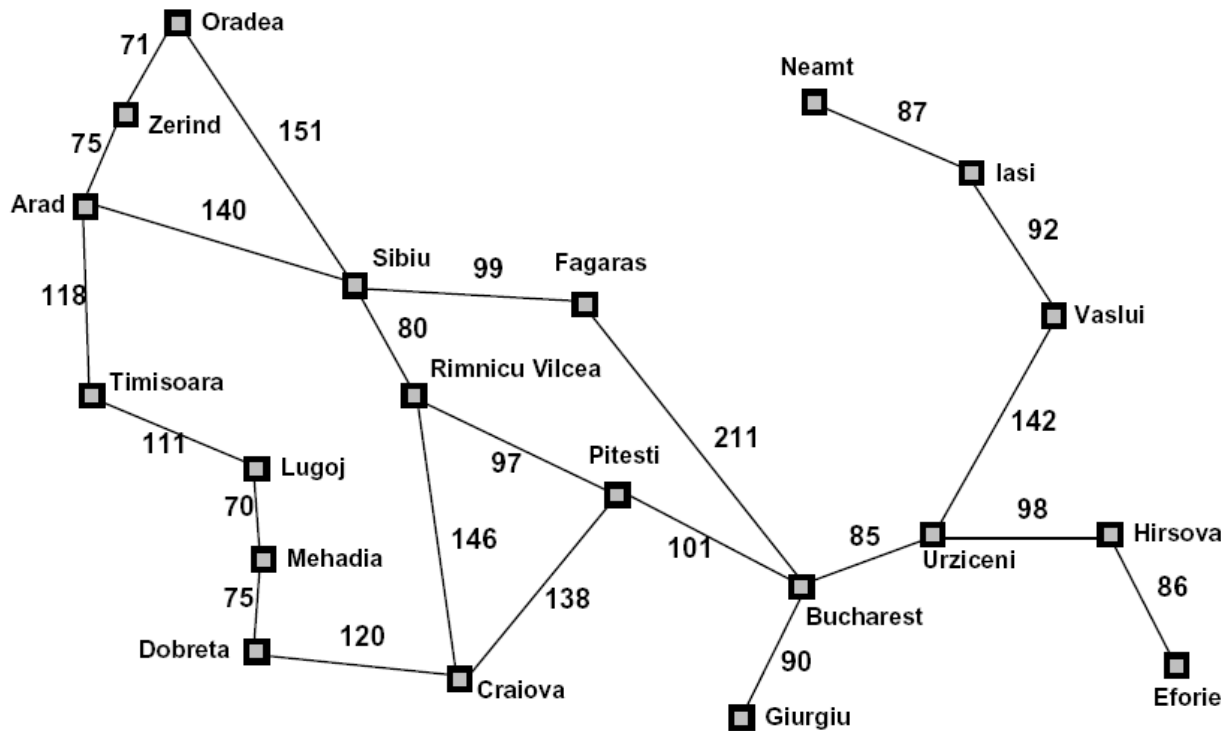
Explored (set)	Frontier (priority queue)
{}	A($0+366=366$)
{A}	Z($75+374=449$), S($140+253=393$), T($118+329=447$)



Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

Explored (set)	Frontier (priority queue)
{}	A($0+366=366$)
{A}	Z($75+374=449$), S($140+253=393$), T($118+329=447$)
{A, S}	Z(449), T(447), O ($140+151+380 = 671$), F ($140+99+178 = 417$), RV ($140+80+193 = 413$)

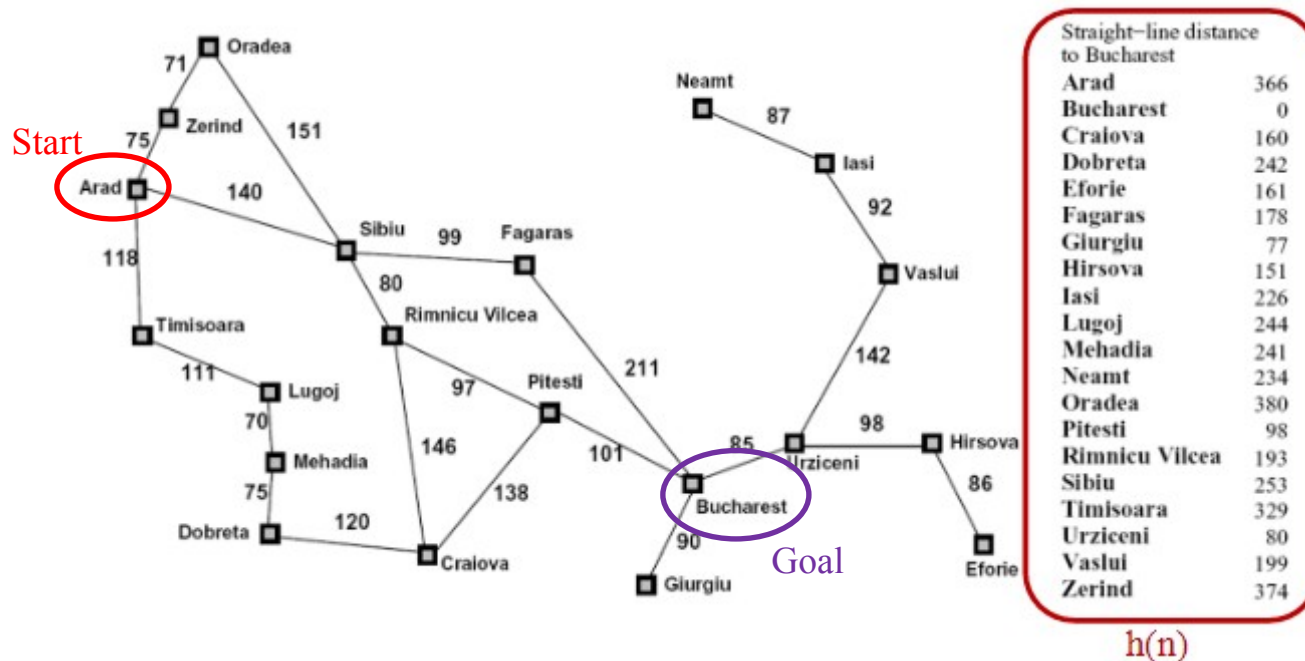


Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374

In-class exercise


- Work on the rest of the iterations for going from Arad to Bucharest
 - Show frontier (priority queue) and explored node



A* SEARCH

PROPERTIES OF A*

Is A* Optimal?



- What went wrong?
- Estimated goal cost $>$ actual goal cost
- We need estimates to be less than actual costs!

Conditions for optimality of A^*


1. Heuristic must be *admissible*
 - Must *never overestimate* true cost
 - Admissible (optimistic) heuristics slow down bad plans but never outweigh true costs

Admissible Heuristics

- A heuristic h is *admissible* (optimistic) if:

where c is the true cost to nearest goal


- What is an acceptable heuristic for node A?



Conditions for optimality of A^*

2. Heuristic must be *consistent* (also called *monotonic*)


- The estimated cost of getting to a goal is **never more** than cost of an action + estimated cost of next state



- Every consistent heuristic is also admissible
- In general, most natural admissible heuristics tend to be consistent

Optimality of A^* Search

- Complete
 - Terminates even on infinite state spaces if a reachable goal exists
- Optimal
 - Always finds the closest goal
- Optimally efficient
 - No other optimal algorithm with the same heuristics expands fewer nodes
- Note: Uniform cost search is a special case of A^* ($h = 0$)




Complexity of A* Search


- Time complexity:
 - Unfortunately, number of states still exponential in the length of the solution
 - Number of states depends on the “error” in the heuristic
 - Error =
- Space complexity:
 - Keeps all generated nodes in memory, so exponential

UCS vs A* Contours

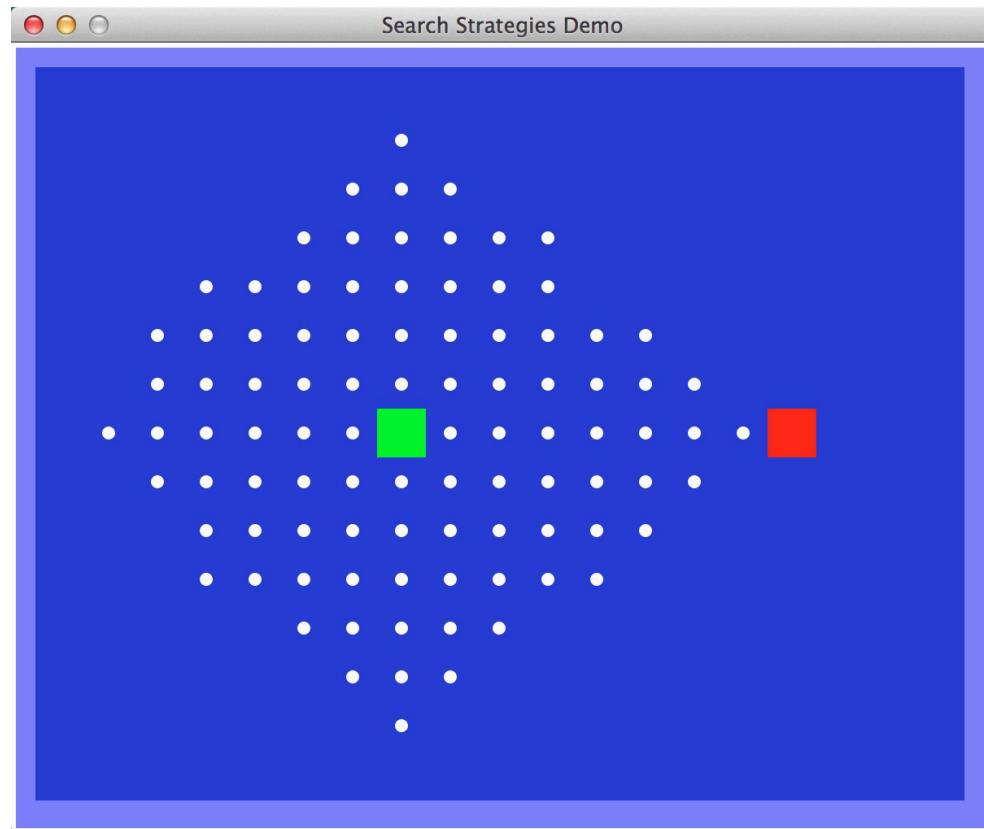
- Uniform-cost expands equally in all “directions”




- A* expands mainly toward the goal, but does hedge its bets to ensure optimality




Video of Demo Contours (Empty) -- UCS



Video of Demo Contours (Empty) -- Greedy



Video of Demo Contours (Empty) – A^*



A* Applications

- Video games
- Path finding
- Robot motion planning
- Resource allocation



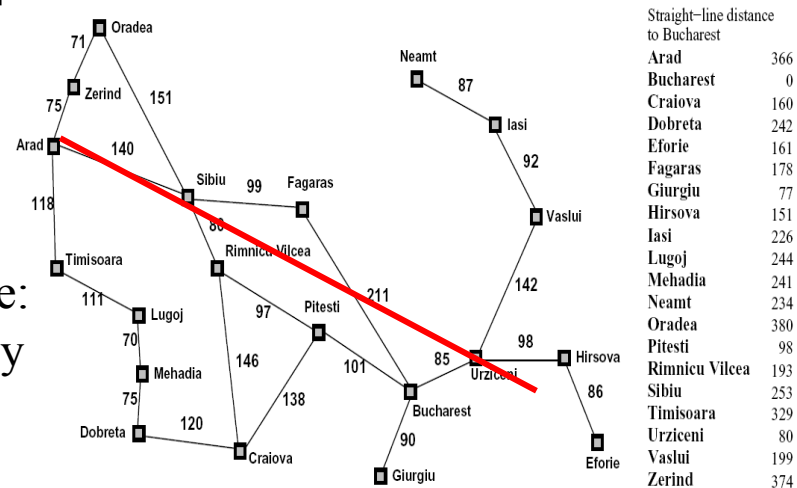
A* SEARCH

CREATING HEURISTICS

Creating Admissible Heuristics

- Most of the work in solving hard search problems optimally is in coming up with admissible heuristics
- Often, admissible heuristics are solutions to *relaxed problems* where new actions are available

Euclidean distance:
assume you can fly




- Inadmissible heuristics are often useful too

Example: 8 Puzzle

7	2	4
5		6
8	3	1

Start State



Actions

	1	2
3	4	5
6	7	8

Goal State

- What are the states?
- How many states?
- What are the actions?
- How many successors from the start state?
- What should the costs be?

8 Puzzle I


- Heuristic: Number of tiles misplaced
- $h(\text{start}) = 8$
- This is a *relaxed-problem* heuristic
- Why is it admissible?

7	2	4
5		6
8	3	1

Start State


	1	2
3	4	5
6	7	8

Goal State



Sliding tile puzzle: misplaced tiles heuristic

5	2	7
8	4	
1	3	6




1	2	3
4	5	6
7	8	

$$h(n) = ?$$

A* : Summary

- A* uses both backward costs and (estimates of) forward costs
- A* is optimal with admissible / consistent heuristics
- Heuristic design is key: often use relaxed problems



Acknowledgement

- <https://inst.eecs.berkeley.edu/~cs188/su20/>

References

- Russel and Norvig, Artificial Intelligence: A Modern Approach, 4th edition, Prentice Hall, 2010.